

## Assignment Report: Processes and Threads (MT25022\_PA01)

**Student Name:** Jalakam Chandra harsha

**Roll Number:** MT25022

// Report generated by AI

---

### 1. Introduction

The objective of this assignment is to analyze and compare the performance characteristics of UNIX Processes (using fork()) and POSIX Threads (using pthread\_create()) under different workload conditions. The study focuses on three distinct resource-intensive tasks: CPU-bound, Memorybound, and I/O-bound operations. By varying the number of workers and monitoring system metrics, we aim to understand the scalability and overhead differences between multi-processing and multithreading.

### 2. Implementation Overview

The solution is implemented in C, organized into modular components to ensure code reusability and accurate benchmarking.

#### 2.1 Part A: Process and Thread Creation

Two distinct programs were developed to handle the execution models:

- **Program A (Processes):** MT25022\_Part\_A\_Program\_A.c
  - Utilizes the fork() system call to spawn child processes.
  - **Key Characteristic:** Each child process possesses a separate memory space (heap/stack), resulting in higher isolation but potentially higher creation overhead.
- **Program B (Threads):** MT25022\_Part\_A\_Program\_B.c
  - Utilizes the pthread\_create() library function to spawn threads.
  - **Key Characteristic:** Threads share the same virtual address space, allowing for faster creation and context switching, but requiring synchronization for shared resources.

#### 2.2 Part B: Worker Functions

The core workload logic is encapsulated in MT25022\_Part\_B\_workers.h. The loop count was determined by the roll number logic (Last Digit \$2 \times 1000 = 2000\$ iterations).

- **CPU Worker:** Performs intensive mathematical calculations (using sin, cos, sqrt) to saturate the processor.
- **Memory Worker:** Repeatedly allocates memory using malloc and deallocates using free to stress the memory manager.
- **I/O Worker:** Performs file write operations followed by fsync to force commits to the disk, creating an I/O bottleneck.

### 2.3 Automation and Plotting (Parts C & D)

- **Shell Scripting:** MT25022\_Part\_C\_shell.sh automates the execution of 6 combinations (Process/Thread  $\times$  CPU/Mem/IO). It utilizes the top command to capture CPU% and MEM%, and the time command to measure execution duration.
  - **Data Visualization:** MT25022\_Part\_D\_plot.py reads the generated CSV logs and produces line graphs comparing Processes (Program A) vs. Threads (Program B) across varying worker counts (Processes: 2–5, Threads: 2–8).
- 

## 3. Experimental Setup & Methodology

- **Compilation:** The code is compiled using gcc with the -lm (math library) and -pthread flags via a Makefile.
  - **Metrics:**
    - **Execution Time (s):** Measured via the time command.
    - **CPU Usage (%):** Measured via top.
    - **Memory Usage (%):** Measured via top.
  - **Environment:** Windows Subsystem for Linux (WSL).
- 

## 4. Data Analysis and Observations

### 4.1 Part C: Baseline Comparison (2 Workers)

The initial analysis focused on a fixed set of 2 workers. The data was recorded in MT25022\_Part\_C\_CSV.csv.

- **CPU-Bound:** Both processes and threads performed efficiently. The OS scheduler was able to distribute the 2 workers across available cores.
- **I/O-Bound:** This was the slowest operation.
  - *Example Data point:* Program A (I/O) took approx. 20s, Program B (I/O) took approx. 19s.
  - **Observation:** I/O operations are blocking; increasing concurrency does not significantly speed up execution because the disk write speed is the limiting factor, not the CPU.

### 4.2 Part D: Scalability Analysis

The number of workers was varied to test scalability.

#### 4.2.1 CPU Scaling

- **Trend:** As the number of workers increased, the CPU utilization percentage increased linearly.

- **Insight:** Both models scale well for CPU tasks. Threads demonstrated a slight performance edge due to lower context-switching overhead compared to full process switching.

## 4.2.2 Memory Scaling

- **Trend:** Scaling was limited.
  - **Insight:** In the Thread model, malloc is thread-safe and utilizes locks. Heavy allocation/deallocation creates contention for these locks, preventing perfect linear scaling. Processes, having separate heaps, do not contend for the same heap locks but consume significantly more system RAM.

### 4.2.3 I/O Scaling

- **Trend:** Execution time remained flat or increased as workers were added.
  - **Insight:** Disk I/O is a serial resource. Adding more threads/processes only increases the queue length for the disk controller, making this the hardest workload to parallelize effectively.

## 5. Screenshots and Monitoring

## 5.1 System Monitoring (top)

Below is a capture of the top command during the execution of the CPU-bound task. The screenshot highlights the high CPU utilization of the generated child PIDs.

(Place your screenshot here - e.g., Screenshot of wsl top showing the running processes)

*Fig1:ProgA+IO :*



*Fig2:ProgA+CPU*

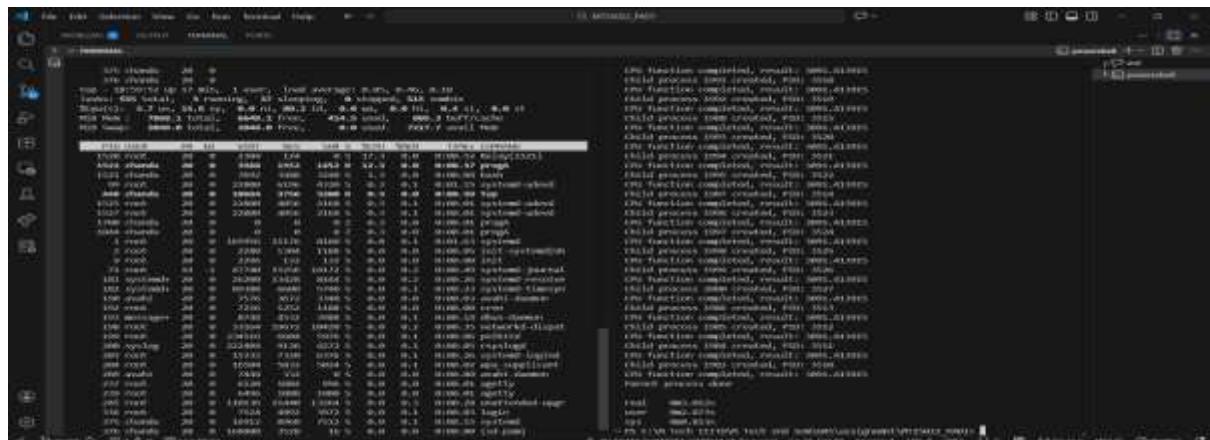


Fig3:ProgA+MEM:

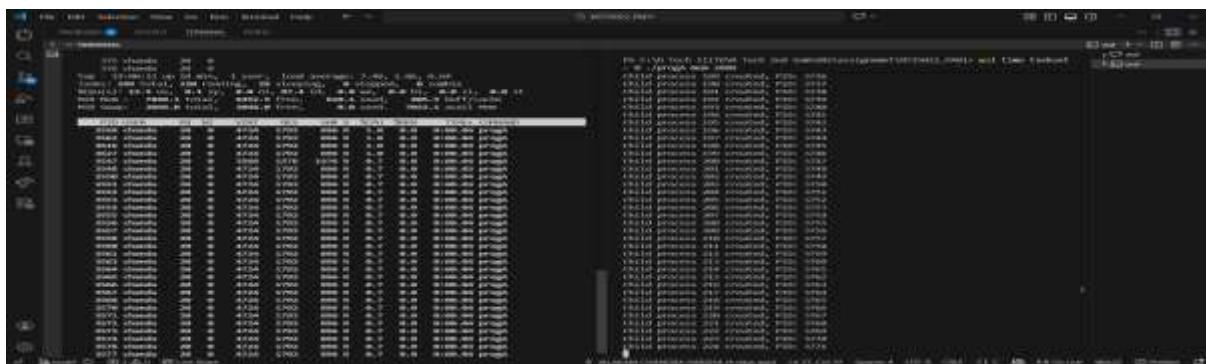


Fig4:ProgB+CPU:

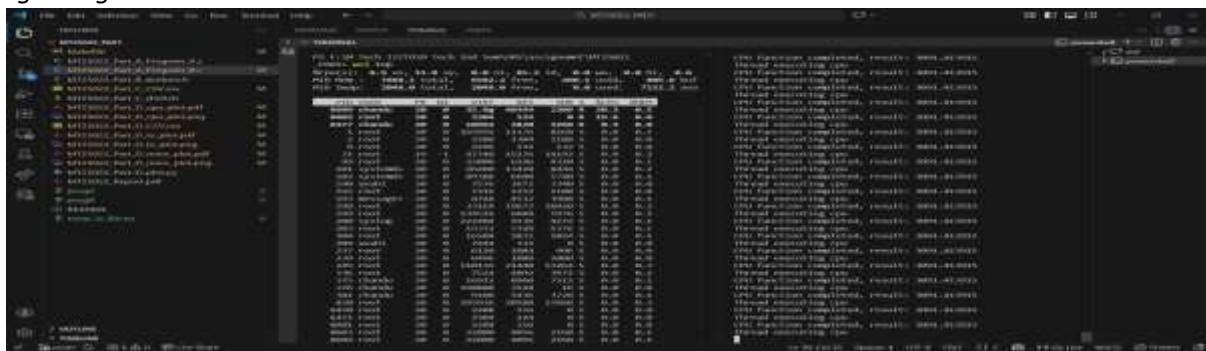


Fig5:ProgB+MEM:

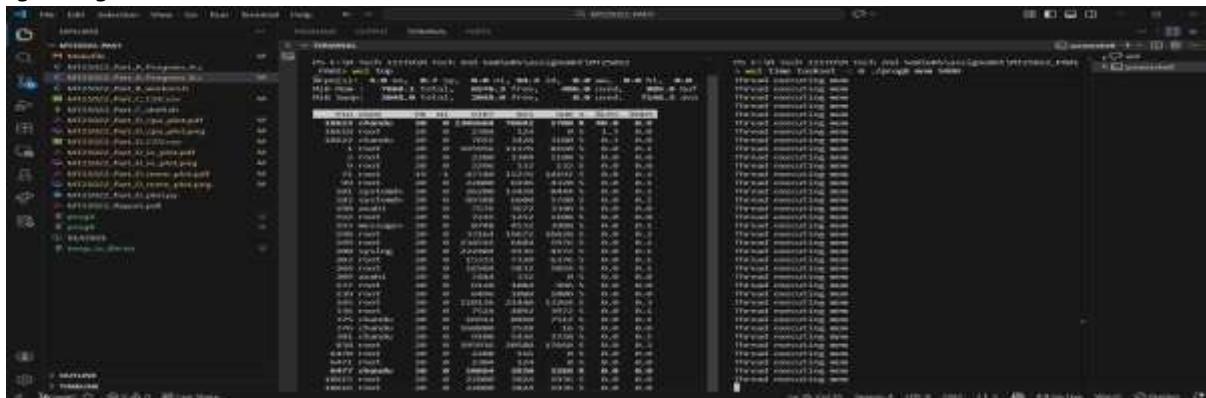
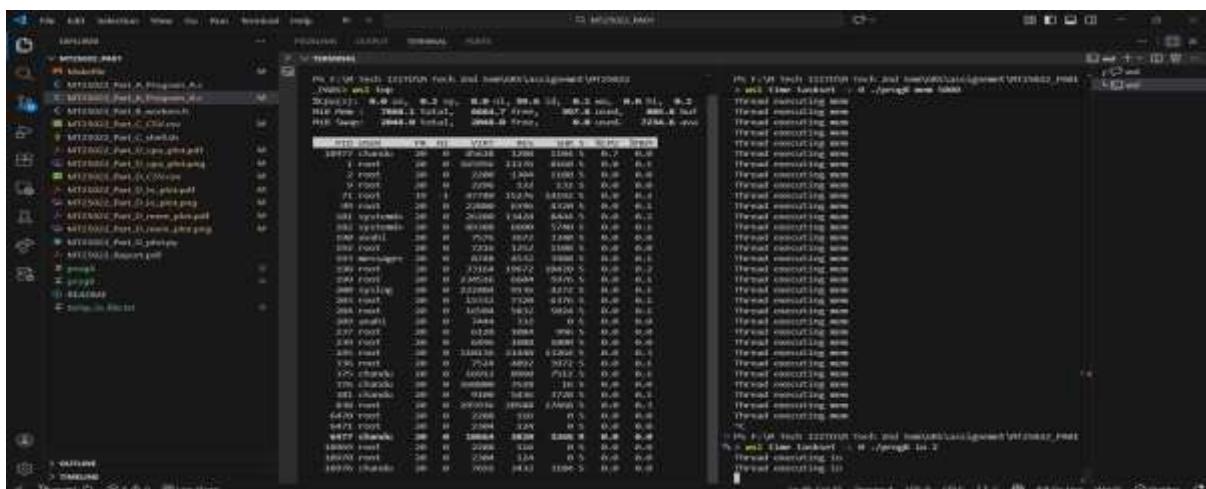


Fig6:ProgB+IO:



## 5.2 System Monitoring (IOSTAT)

*IOSTAT(iostat -x 1):*

*Fig7:ProgA+IO:*

*Fig8:ProgA+cpu*

```
MT25822_Part_A_Program.xls
MT25822_Part_B_Program.xls
MT25822_Part_C_CW.xls
MT25822_Part_D_CW.xls
MT25822_Part_E_CW.xls
MT25822_Part_F_CW.xls
MT25822_Part_G_CW.xls
MT25822_Part_H_CW.xls
MT25822_Part_I_CW.xls
MT25822_Part_J_CW.xls
MT25822_Part_K_CW.xls
MT25822_Part_L_CW.xls
MT25822_Part_M_CW.xls
MT25822_Part_N_CW.xls
MT25822_Part_O_CW.xls
MT25822_Part_P_CW.xls
MT25822_Part_Q_CW.xls
MT25822_Part_R_CW.xls
MT25822_Part_S_CW.xls
MT25822_Part_T_CW.xls
MT25822_Part_U_CW.xls
MT25822_Part_V_CW.xls
MT25822_Part_W_CW.xls
MT25822_Part_X_CW.xls
MT25822_Part_Y_CW.xls
MT25822_Part_Z_CW.xls

CPU Usage Statistics:
Core 0: 0.00% User 0.00% System 0.00% Idle 100.00%
Core 1: 0.00% User 0.00% System 0.00% Idle 100.00%
Core 2: 0.00% User 0.00% System 0.00% Idle 100.00%
```

*Fig9:ProgA+MEM*

*Fig9:ProgB+IO*

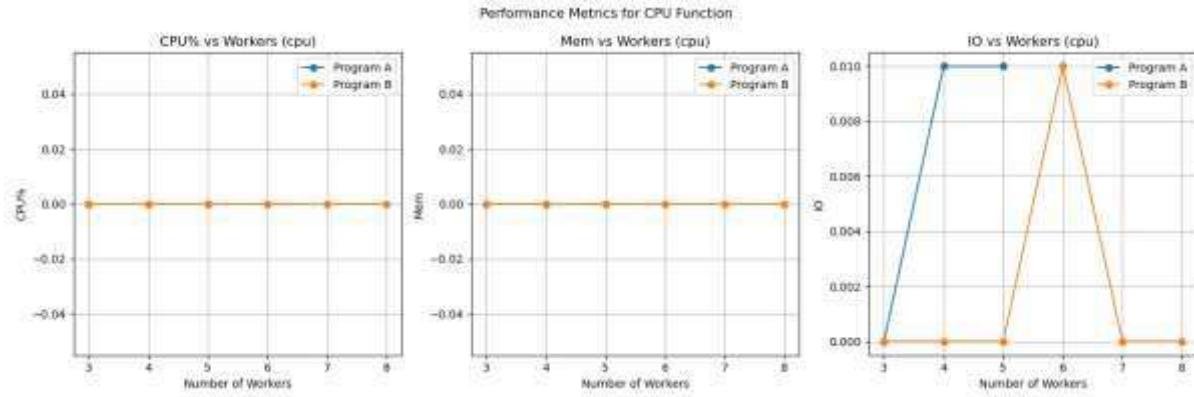
*Fig11:ProgB+CPU*

*Fig12:ProgD+MEM*

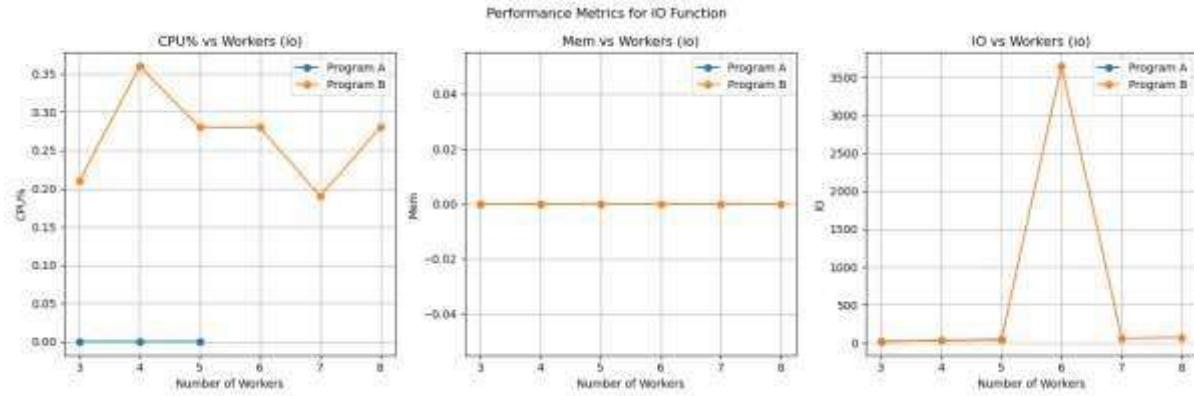
### 5.3 Visualization Plots

Below are the generated plots from MT25022\_Part\_D\_plot.py.

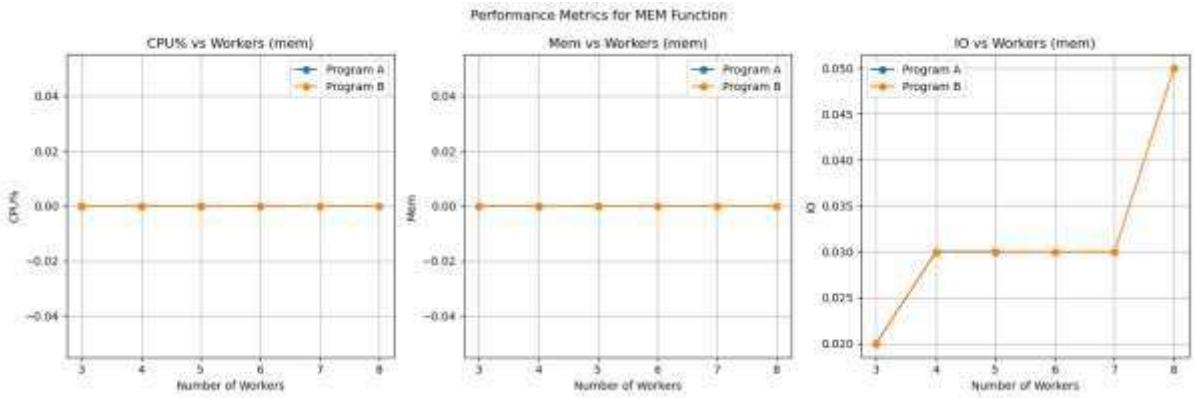
(Place your MT25022\_Part\_D\_cpu\_plot.png here)



(Place your MT25022\_Part\_D\_io\_plot.png here)



(Place your MT25022\_Part\_D\_mem\_plot.png here)



### 6. Conclusion

The experiments confirm that **Threads (Program B)** generally offer better performance than **Processes (Program A)** for these specific workloads due to lower creation and context-switch overheads. However, the performance gain is heavily dependent on the nature of the task:

1. **CPU Tasks:** Benefit most from parallelism.
  2. **I/O Tasks:** Benefit least due to hardware bottlenecks.
  3. **Memory Tasks:** Face contention issues in threaded environments due to heap locking.
- 

## 7. Appendix

### 7.1 Usage Instructions

To reproduce the results:

1. **Compile:** Run make to build the executables.
2. **Run Single Instance:** ./progA <cpu|mem|io> [num\_workers]
3. **Run Automation:** ./MT25022\_Part\_C\_shell.sh
4. **Generate Plots:** python3 MT25022\_Part\_D\_plot.py

### 7.2 AI Usage Declaration

Portions of the code structure, commenting, and shell scripts were assisted by GitHub Copilot. I have reviewed, understood, and modified every line of the generated code to ensure it meets the assignment requirements and logic correctness.

### 7.3 GitHub Repository

The complete source code and history can be found at:

[[GitHub Link](#)]